

Deploy to Kubernetes

Estimated time needed: 60 minutes

Welcome to the **Deploy to Kubernetes** hands-on lab. Now that your microservice is built and tested, it is time to deploy it to a Kubernetes environment to run it. In particular, you will use OpenShift, which is based on Kubernetes and adds additional developer capabilities.

Objectives

In this lab, you will:

- Take the next story from the Sprint Backlog to work on
- Create a Dockerfile and build an image from your microservice
- Create Kubernetes manifests for your deployment
- Deploy your Docker image in an OpenShift Kubernetes cluster
- View the logs to ensure your service is running
- Make a pull request and merge your changes
- Move the story to Done

Note: Important Security Information

Welcome to the Cloud IDE with OpenShift. This is where all your development will take place. It has all the tools you will need to use Docker for deploying a PostgreSQL database.

It is important to understand that the lab environment is **ephemeral**. It only lives for a short while before it is destroyed. It is imperative that you push all changes made to your own GitHub repository so that it can be recreated in a new lab environment any time it is needed.

Also note that this environment is shared and therefore not secure. You should not store any personal information, usernames, passwords, or access tokens in this environment for any purposes.

Your Task

1. If you haven't generated a **GitHub Personal Access Token** you should do so now. You will need it to push code back to your repository. It should have `repo` and `write` permissions, and be set to expire in `60` days. When Git prompts you for a password in the Cloud IDE environment, use your Personal Access Token instead.
2. The environment may be recreated at any time so you may find that you have to perform the **Initialize Development Environment** each time the environment is created.

Note on Screenshots

Throughout this lab, you will be prompted to take screenshots and save them on your device. You will need these screenshots to either answer graded quiz questions or upload as your submission for peer review at the end of this course. Your screenshot must have either the `.jpg` or `.png` extension.

To take screenshots, you can use various free screen-capture tools or your operating system's shortcut keys. For example:

- **Mac:** you can use `Shift + Command + 3` (`⌘` + `⇧` + `3`) on your keyboard to capture your entire screen, or `Shift + Command + 4` (`⌘` + `⇧` + `4`) to capture a window or area. They will be saved as a file on your Desktop.
- **Windows:** you can capture your active window by pressing `Alt + Print` Screen on your keyboard. This command copies an image of your active window to the clipboard. Next, open an image editor, paste the image from your clipboard to the image editor, and save the image.

Initialize Development Environment

Because the Cloud IDE with OpenShift environment is ephemeral, it may be deleted at any time. The next time you come into the lab, a new environment may be created. Unfortunately, this means that you will need to initialize your development environment every time it is recreated. This shouldn't happen too often as the environment can last for several days at a time but when it is removed, this is the procedure to recreate it.

Overview

Each time you need to set up your lab development environment you will need to run three commands.

Each command will be explained in further detail, one at a time, in the following section.

{your_github_account} represents your GitHub account username.

The commands include:

```
1. 1
2. 2
3. 3
4. 4

1. git clone https://github.com/{your_github_account}/devops-capstone-project.git
2. cd devops-capstone-project
3. bash ./bin/setup.sh
4. exit
```

Copied!

Now, let's discuss each of these commands and explain what needs to be done.

Task Details

Initialize your environment using the following steps:

1. Open a terminal with `Terminal -> New Terminal` if one is not open already.
2. Next, use the `export GITHUB_ACCOUNT=` command to export an environment variable that contains the name of your GitHub account.

Note: Substitute your real GitHub account for the {your_github_account} placeholder below:

```
1. 1
1. export GITHUB_ACCOUNT={your_github_account}
```

3. Then use the following commands to clone your repository, change into the `devops-capstone-project` directory, and execute the `./bin/setup.sh` command.

```
1. 1
2. 2
3. 3

1. git clone https://github.com/{GITHUB_ACCOUNT}/devops-capstone-project.git
2. cd devops-capstone-project
3. bash ./bin/setup.sh
```

Copied! Executed!

You should see the follow at the end of the setup execution:

```
*****
Capstone Environment Setup Complete
*****

Use 'exit' to close this terminal and open a new one to initialize the environment

/home/project/devops-capstone-project$
```

4. Finally, use the `exit` command to close the current terminal. The environment will not be fully active until you open a new terminal in the next step.

```
1. 1
1. exit
```

Copied! Executed!

Validate

In order to validate that your environment is working correctly, you must open a new terminal because the Python virtual environment will only activate when a new terminal is created. You should have ended the previous task by using the `exit` command to exit the terminal.

1. Open a terminal with `Terminal -> New Terminal` and check that everything worked correctly by using the `which python` command:

Your prompt should look like this:

```
(venv) theia:project$
```

Check which Python you are using:

```
1. 1
1. which python
```

Copied! Executed!

You should get back:

```
(venv) theia:project$ which python
/home/theia/venv/bin/python
(venv) theia:project$
```

Check the Python version:

```
1. 1
1. python --version
```

[Copied!](#) [Executed!](#)

You should get back some patch level of Python 3.9:

```
(venv) theia:project$ python --version
Python 3.9.15
(venv) theia:project$
```

This completes the setup of the development environment. Anytime your environment is recreated, you will need to follow this procedure.

You are now ready to start working.

Exercise 1: Pick Up the First Story

The first thing you need to do is to go to your Zenhub kanban board in GitHub to get a story to work on. Take the first story from the top of the Sprint Backlog, move it to In Progress, assign it to yourself, and read the contents.

Your Task

1. Go to your Zenhub kanban board and take the first story from the top of the Sprint Backlog. It should be titled: "Containerize your microservice using Docker".
2. Move the story to In Progress.
3. Open the story and assign it to *yourself*.
4. Read the contents of the story.

Results

The story should look similar to this:

Containerize your microservice using Docker

As a developer

I need to containerize my microservice using Docker

So that I can deploy it easily with all of its dependencies

Assumptions

- Create a Dockerfile for repeatable builds
- Use a Python:3.9-slim image as the base
- It must install all of the Python requirements
- It should not run as root
- It should use the gunicorn wsgi server as an entry point

Acceptance Criteria

1. 1
 2. 2
 3. 3
1. Given the Docker image named accounts has been created
 2. When I use 'docker run accounts'
 3. Then I should see the accounts service running in Docker

[Copied!](#)

You are now ready to begin working on your story.

Exercise 2: Create a Dockerfile

In reading your story you see that the assumptions state that you must create a Dockerfile with the following attributes:

- Create a Dockerfile for repeatable builds
- Use a Python:3.9-slim image as the base
- It must install all of the Python requirements
- It should not run as root
- It should use the gunicorn wsgi server as an entry point

Let's take these in order.

Your Task

1. Change to your project directory: `cd devops-capstone-project`.
2. Use the `git checkout -b add-docker` command to create a new branch called `add-docker` to work on in the development environment.
3. Run `nosetests` and make sure that all of the test cases are passing. Fix any failing tests before proceeding.
4. In the root of the repository, create a file named: `Dockerfile`.
5. Edit the `Dockerfile` and start it from the `python:3.9-slim` image.
 - ▶ [Click here for the answer.](#)
6. Establish a `WORKDIR` of `/app`, copy the `requirements.txt` file into the working directory in the image, and run the `pip` command to install the requirements using the `--no-cache-dir` option to keep the image small.
 - ▶ [Click here for the answer.](#)
7. Copy the `service` package into the working directory of the same name in the image.
 - ▶ [Click here for the answer.](#)
8. Create a non-root user called `theia`, change the ownership of the `/app` folder recursively to `theia`, and switch to the `theia` user.
 - ▶ [Click here for the answer.](#)
9. Finally, `EXPOSE` port `8080` and create a `CMD` statement that runs: `gunicorn --bind=0.0.0.0:8080 --log-level=info service:app`
 - ▶ [Click here for the answer.](#)

Results

You can check that your `Dockerfile` looks like the following:

▶ [Click here to check your work.](#)

Exercise 3: Create a Docker Image

Now that you have created a `Dockerfile`, it's time to create an image from it to see if it works.

Your Task

1. Open a terminal and use the `docker build` command to build a Docker image called `accounts` from the `Dockerfile`.
 - ▶ [Click here for the answer.](#)
2. Use the `docker run` command to test that your image works properly. The PostgreSQL database is running in a Docker container named `postgres` so you will need to `--link postgres` and set the environment variable `DATABASE_URI` to point to it. You might also want to use the `--rm` flag to remove the container when it exists.

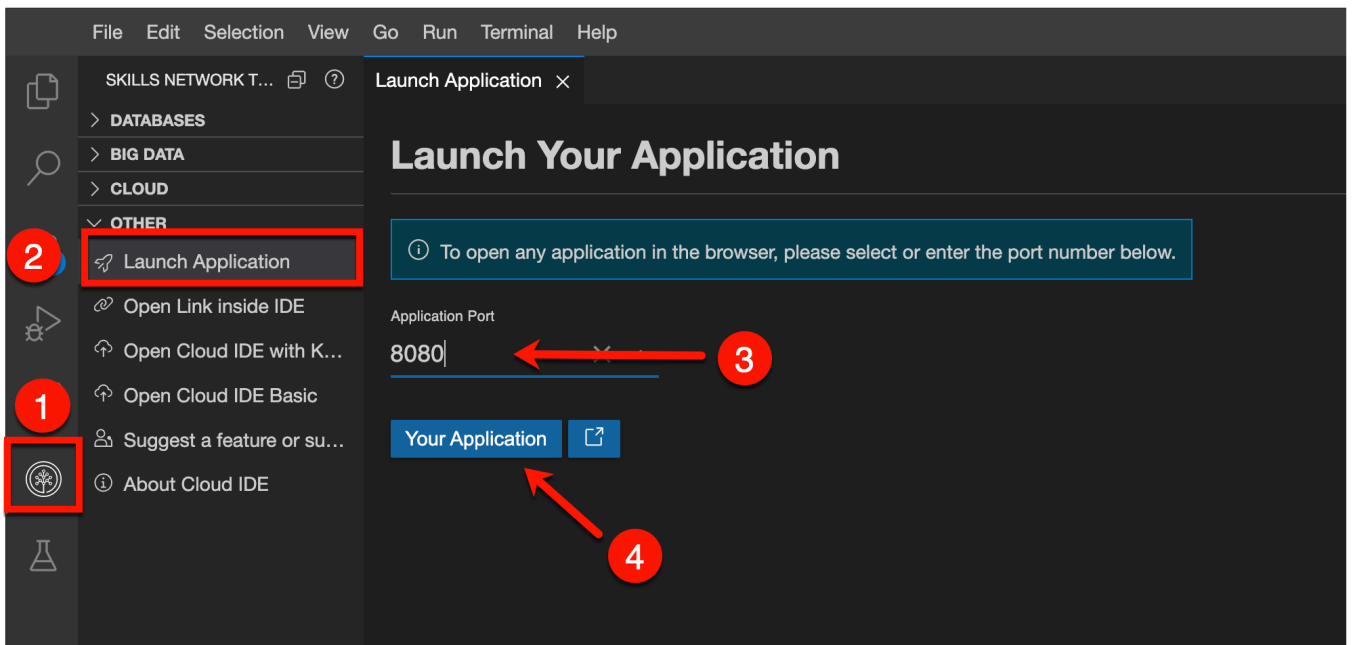
If it worked, you should see the message:

```
1. 1
1. ... [INFO] [__init__] Service initialized!
```

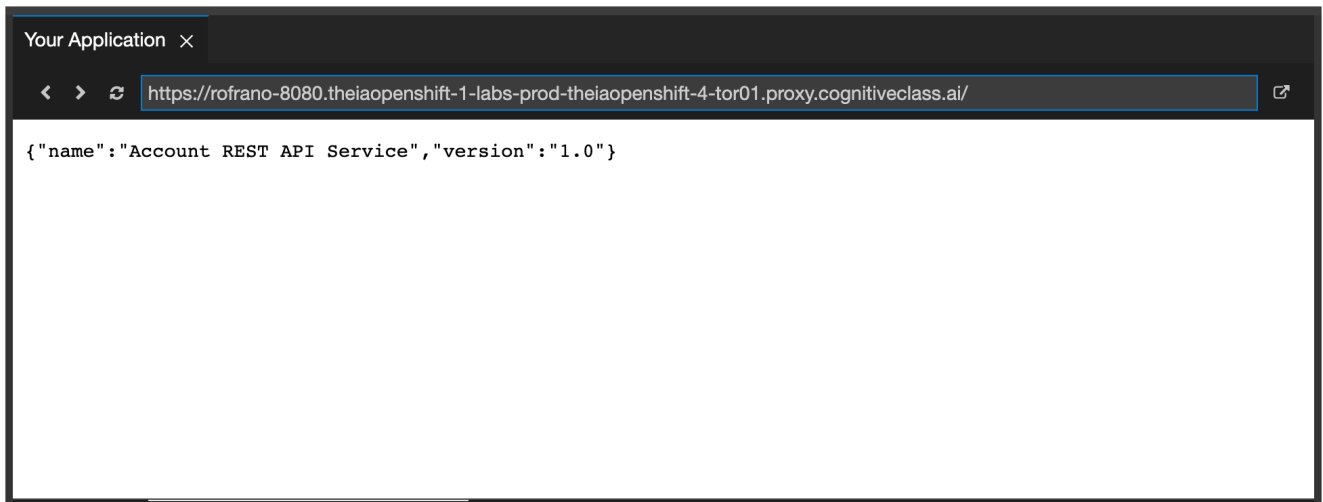
[Copied!](#)

▶ [Click here for the answer.](#)

3. Check that your application is running by (1) clicking the `Skills Network` icon, (2) selecting `Other -> Launch Application`, (3) entering an Application Port of `8080`, and (4) clicking the `Your Application` button.



You should see the following:



4. Use `Ctrl+C` to stop your container.
5. Tag the image as `us.icr.io/$SN_ICR_NAMESPACE/accounts:1` and push it to the IBM Cloud registry.
Note: The environment variable `SN_ICR_NAMESPACE` contains your image namespace in the IBM Cloud Container Registry.
► [Click here for the answer.](#)

Evidence

For the evidence, take a screenshot of the internal web browser's output from **task 3**.

1. Open your application in the browser and save a screenshot of the page as `kube-app-output.jpg` (or `kube-kube-app-output.png`).

Exercise 4: Make a Pull Request

Now that you have a working Docker image, it's time to push the Dockerfile up to GitHub and make a pull request, merge the request, and move your story to Done.

Your Task

1. Use `git status` to make sure that you have committed your changes locally in the development environment.
 2. Use the `git add` command to add the new `Dockerfile` to the staging area.
 3. Commit your changes using the message `Added docker support`.
 4. Push your local changes to a remote branch.
- Note: Use your GitHub **Personal Access Token** as your password in the Cloud IDE environment. You may also need to configure Git the first time you use it with:

```
1. 1
2. 2
1. git config --global user.email "you@example.com"
2. git config --global user.name "Your Name"
```

[Copied!](#)

► [Click here for the answer.](#)

5. Make a pull request, which should kick off the GitHub Actions that are now enabled on your repository.
6. Once the test cases pass, merge your pull request.
7. Move your story to the Done column on your kanban board.
8. Pull the last code down to your development environment and delete your old branch.

```
1. 1
2. 2
3. 3
1. git checkout main
2. git pull
3. git branch -d add-docker
```

[Copied!](#)

Evidence

For the evidence, take a screenshot of your kanban board to show the story is done.

1. Open your kanban board and save a screenshot of the board with your story in the Done column as `kube-docker-done.jpg` (or `kube-docker--done.png`).

Exercise 5: Pick Up the Next Story

It's now time to go to your kanban board in GitHub to get the next story to work on. It should be at the top of the Sprint Backlog.

Your Task

1. Go to your kanban board and take the next story from the top of the Sprint Backlog. It should be titled: "Deploy your Docker image to Kubernetes".
2. Move the story to In Progress.
3. Open the story and assign it to yourself.
4. Read the contents of the story.

Results

The story should look similar to this:

Deploy your Docker image to Kubernetes

As a service provider
I need my service to run on Kubernetes
So that I can easily scale and manage the service

Assumptions

- Kubernetes manifests will be created in yaml format
- These manifests could be useful to create a CD pipeline
- The actual deployment will be to OpenShift

Acceptance Criteria

1. 1
 2. 2
 3. 3
1. Given the Kubernetes manifests have been created
 2. When I use the oc command to apply the manifests
 3. Then the service should be deployed and run in Kubernetes

[Copied!](#)

You are now ready to begin working on your second story.

Exercise 6: Deploy to Kubernetes

For the "Deploy to Kubernetes" story you must create the manifests required to consistently deploy your microservice. At some point in the future, you need to create a CD pipeline to perform continuous delivery so, while you are deploying manually now, it's important that you create manifests that can be used later in the pipeline.

You are going to need a PostgreSQL database in Kubernetes for your application to use. Luckily, you are using OpenShift, which comes with a number of templates for creating services. Your first task is to deploy the postgresql-ephemeral template, which will create an ephemeral PostgreSQL database with no backing storage for test purposes.

Your Task

1. Use the `git checkout -b add-kubernetes` command to create a new branch called `add-kubernetes` to work on in the development environment.
2. Use the `oc new-app` command to deploy the `postgresql-ephemeral` template.
 - ▶ [Click here for the answer.](#)
3. Use `oc get pods` to make sure that the postgres service is defined and the postgres pod is running.

Results

```
(venv) theia:devops-capstone-project$ oc get all
NAME                                READY    STATUS    RESTARTS    AGE
pod/openshift-web-console-8bd9fcbf8-2mlw8    2/2      Running    0            2d2h
pod/openshift-web-console-8bd9fcbf8-d8wts    2/2      Running    0            2d2h
pod/postgresql-1-deploy                0/1      Completed    0            177m
pod/postgresql-1-p7rfz                1/1      Running    1 (11/6m ago)  177m

NAME                                DESIRED    CURRENT    READY    AGE
replicationcontroller/postgresql-1    1          1          1        177m

NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/openshift-web-console    ClusterIP    172.17.114.10    <none>          8080/TCP    2d2h
service/postgresql                ClusterIP    172.21.42.231    <none>          5432/TCP    177m

NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/openshift-web-console    2/2      2             2            2d2h

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/openshift-web-console-8bd9fcbf8    2          2          2        2d2h

NAME                                REVISION    DESIRED    CURRENT    TRIGGERED BY
deploymentconfig.apps.openshift.io/postgresql    1           1          1          config,image(postgresql:10-el8)
(venv) theia:devops-capstone-project$
```

You are now ready to create Kubernetes manifests for your microservice.

Exercise 7: Create Manifests

Here is a tip for getting started creating manifest yaml files. You can use the `kubectl` or `oc` CLI to create a deployment or service and capture the definition in a yaml file by adding the flags `--dry-run=client -o yaml`. This code doesn't actually create anything (`--dry-run=client`) but sends output to yaml (`-o yaml`). Then all you need to do is redirect that to a file.

1. Create a manifest definition for the account deployment using the `oc create deployment` command with the `--dry-run -o yaml` option and redirect it to a file called `deploy/deployment.yaml`. Specify the image that you pushed to the IBM Cloud registry and request three replicas.
 - ▶ [Click here for the answer.](#)
2. Your microservice needs to know the details about the postgres database that you just deployed. In particular, it needs the following environment variables: `DATABASE_HOST`, `DATABASE_NAME`, `DATABASE_PASSWORD`, and `DATABASE_USER`. Use the `oc describe` command to see what keys are in the secret that you can use:
 - ▶ [Click here for the answer.](#)
3. Edit the `deploy/deployment.yaml` file and use the keys that you found in the secret along with a `DATABASE_HOST` of `postgresql` to add the required environment variables to the manifest.
 - ▶ [Click here for a hint.](#)
 - ▶ [Click here for the answer.](#)
4. Apply this deployment using `oc create` and point it to your `deploy/deployment.yaml` file.
 - ▶ [Click here for the answer.](#)
5. Create a manifest definition for the account service using the `oc expose` command, using a type of `NodePort`, and a Port of `8080`, and redirect it to a file called `deploy/service.yaml`.
 - ▶ [Click here for the answer.](#)
6. Apply this service using the `oc create` command and point it to your `deploy/service.yaml` file.
 - ▶ [Click here for the answer.](#)
7. Now it's time to see if everything is running. Use `oc get all` and filter by the level `app=accounts` to see your deployment running.
 - ▶ [Click here for the answer.](#)

You should see something similar to the following:

```
(venv) theia:devops-capstone-project$ oc get all -l app=accounts
NAME                                READY    STATUS    RESTARTS    AGE
pod/accounts-7f4df674b9-dhm49      1/1      Running   0            3m58s

NAME            TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/accounts ClusterIP    172.21.183.7  <none>         8080/TCP   23s

NAME            READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/accounts 1/1      1              1            3m58s

NAME            DESIRED    CURRENT    READY    AGE
replicaset.apps/accounts-7f4df674b9 1          1          1        3m58s
(venv) theia:devops-capstone-project$
```

Note: There should be a deployment, replicaset, pod, and service.

8. Finally, expose your service using an OpenShift route. Use the `oc create` command to create a route called `accounts` with edge termination that exposes the `--service` named `accounts`.

► [Click here for the answer.](#)

9. Use the `oc get routes` command to get the route that was assigned to your service.

```
1. 1
1. 1. oc get routes
Copied! Executed!
```

10. Copy the URL of your route and paste it into your browser to see your application running in OpenShift.

Exercise 8: Make Another Pull Request

Now that you have a working deployment, it's time to push the Kubernetes manifests up to GitHub, make a pull request, merge the request, and move your story to done.

Your Task

- 1. Use `git status` to make sure that you have committed your changes locally in the development environment.
- 2. Use the `git add` command to add the new `deployment.yaml` and `service.yaml` to the staging area.
- 3. Commit your changes using the message `Added Kubernetes support`.
- 4. Push your local changes to a remote branch.

Note: Use your GitHub **Personal Access Token** as your password in the Cloud IDE environment. You may also have to configure Git the first time you use it with:

- 5. Make a pull request.
- 6. Once the test cases pass, merge your pull request.
- 7. Move your story to the Done column on your kanban board.
- 8. Pull the latest code down to your development environment and delete your old branch.

```
1. 1
1. 1. 2
1. 2. 3
1. 3. git checkout main
1. 4. git pull
1. 5. git branch -d add-kubernetes
Copied!
```

Evidence

For the evience, take a screenshot of your kanban board to show the story is done.

- 1. Open your kanban board and save a screenshot of the board with your story in the Done column as `kube-kubernetes-done.jpg` (or `kube-kubernetes-done.png`).

Collect Final Evidence

You need to collect the following evidence as proof of completion of this lab.

- 1. Save the URL link to your `Dockerfile` on GitHub. Just open the file on GitHub and save the URL. You will need to provide this when you submit your evidence.
- 2. Issue the command `docker image ls` and save a screenshot of the output as `kube-images.jpg` (or `kube-images.png`).
- 3. Issue the command `oc get all -l app=accounts` and save a screenshot of the output as `kube-deploy-accounts.jpg` (or `kube-deploy-accounts.png`).

Conclusion

Congratulations! You have built a Docker image from a `Dockerfile` and deployed that image to an OpenShift Kubernetes cluster using `yaml` manifests that can be reused in a continuous delivery (CD) pipeline.

Next Steps

Implement the third story in Sprint 3.

Author

[John J. Rofrano](#)
Updated story markdown formatting

Other Contributor(s)

Change Log

Date	Version	Changed by	Change Description
2022-10-14	0.1	John Rofrano	Initial version created
2022-10-25	0.2	Beth Larsen	QA pass
2022-10-28	0.3	John Rofrano	Updated story markdown formatting